




ChatMaJA

From Data to Answers: Developing a Medical Q&A System with LLMs and Transformers

BY : AGATA KACZMAREK, PRANJAL SHARMA, JAN SMOLEŃ, MATEUSZ STĄCZEK

ChatMaJA



CHATMAJA IS A QUESTION &
ANSWERING MACHINE BASED
ON LLMS AND TRANSFORMER.



IT IS BASED ON PAST 10 YEARS
MEDICAL DATA

Data Collection

- ▶ We simply extracted data in multiple excel files containing PMIDs from PubMed
- ▶ These PMIDs were fed into PubMed2XL.com to retrieve title and abstracts.
- ▶ Finally, all the files were merged into single master sheet.

Web UI Designing

- ▶ We used Flask Framework to develop WebUI for this question-answering machine.
- ▶ It is developed locally over the docker



Abstract embedding

- ▶ Abstract are split into chunks, each consisting of 3 sentences
- ▶ Each chunk is embedded separately using BERT from HuggingFace. We use mean pooling to create vectors of even length
- ▶ Embedding are stores in .csv file, which will be then used to populate vector databse
- ▶ For research and demonstration purpose, we only processed first 100 articles

Vector Database - Milvus

- ▶ We are using a vector database Milvus run locally in a Docker container
- ▶ Our collection consists of fields:
 - Id – in the form of {PMID}{chunk_number}
 - title
 - chunk
 - chunk_embedded



Relevant document retrieval

- ▶ Search of relevant chunk is conducted by embedding the query and using cosine similarity
- ▶ Top k relevant chunks are returned. Id, Title of the article and string representation of chunk are also included.

Pros and cons of the approach

- + Milvus offers a very convenient docker image, paired with good quality Python interface and implemented vector similarity search

- + BERT contextual embeddings are relatively simple to compute and work quite well in many cases

- the system struggles with handling keywords

- so far, everything is only tested locally on a small portion of data

- The best way the data is divided into chunks needs to be further researched

Future work - basic

- ▶ Choose LLM – at the beginning we will try one of GPTs
- ▶ Add prompt generation to use LLM to generate easily readable answer
- ▶ Create Test Data – it will probably be based on PubMedQA dataset
- ▶ Integrate with our UI
- ▶ Evaluate

Improving our solution

We are not satisfied with the performance of the current solution, therefore:

- ▶ **Data acquisition – automate it** (now it's a very *manual process*)
- ▶ **Chunking** – investigate chunking methods from LangChain
 - ▶ **Current:** split after every 3 sentences
 - ▶ **New** (example): *RecursiveCharacterTextSplitter*, with overlap
- ▶ Similarity measure – check if changing it matters
- ▶ Questions types – might try more than just yes/no.
- ▶ **LangChain – integrate all components** of our solution **using Chains**

The end – thank you for your attention!
Discussion time

